



Project #1: Artificial Neural Networks

Amir Sadovnik

COSC 525: Deep Learning (Spring 2020)

1 Overview

In this project you will be implementing a simple artificial neural network (ANN) library. I will ask you to implement it in a specific object oriented manner. Although this will make the library less efficient it will allow you to gain a better understanding of how ANN's predict and learn.

Make sure to use your own solution and your own code.

2 Problem Description

You are tasked with writing the following functions:

1. Write a **Neuron** class.
 - (a) Should be initialized with an activation function, the number of inputs, the learning rate, and possibly a vector of weights (if not set to random).
 - (b) Should have an **activate** method which given a value returns its value after activation (depending on the activation function used).
 - (c) Should have a **calculate** method which given an input calculates the output.
 - (d) Any additional methods/instance variables needed for back-propagation
2. Write a **FullyConnectedLayer** class.
 - (a) Should be initialized with the number of neurons in the layer, the activation function for all the neurons in the layer, the number of inputs, the learning rate, and possibly a vector of weights (if not set to random).
 - (b) Should have a **calculate** method which given an input calculates the output of all the neurons in the layer.
 - (c) Any additional methods/instance variables needed for back-propagation

3. Write a `NeuralNetwork` class.

- (a) Should be initialized with the `number of layers`, `number of neurons in each layer`, a `vector of activation functions for each layer`, the `number of inputs to the network`, the `loss function` used, the `learning rate`, and possibly a vector of weights (if not set to random).
 - (b) Should have a `calculate` method which given an input calculates the output of the network.
 - (c) Should have a `calculateloss` method which given an input and desired output calculates the loss.
 - (d) Should have a `train` method, which given a single input and desired output takes one step of gradient descent.
 - (e) Any additional methods/instance variables needed for back-propagation
4. Your `main` method should take the following command line variables to show the following:
- (a) If given `example`, simply run a single step of back-propagation using the example we did in class.
 - (b) If given `and` should train a single perceptron to produce the output of the “and” logic gate. Show the results of the prediction for all 4 possible inputs.
 - (c) If given `xor` should train two types of networks to produce the output of the “xor” logic gate. One is a single perceptron, while the other adds a hidden layer. Show the results of the prediction for all 4 possible inputs for each of them.

3 Additional Information

You must do so under the following constraints:

- 1. Your library should support at least two activation functions: logistic and linear
- 2. Your library should support at least two loss functions: square error and binary cross entropy loss
- 3. You must use Python3 with only the numpy and sys libraries allowed
- 4. Make sure to comment your code

4 Report

You should submit a short PDF report with the following (if you do not have anything to add for a section, simply put the section title and then state there is nothing to add):

- 1. A short introduction to the problem.

2. Assumptions/choices you have made.
3. Problems/Issues you were not able to solve.
4. How to run your code (if there is any difference from what is stated)
5. Show a graph of the loss dropping as a function of learning for different learning rates. Try to find learning rates that are too high/low. Describe what you see in the graph in words.

5 Submission

You are required to submit one zip file with the code and your report.